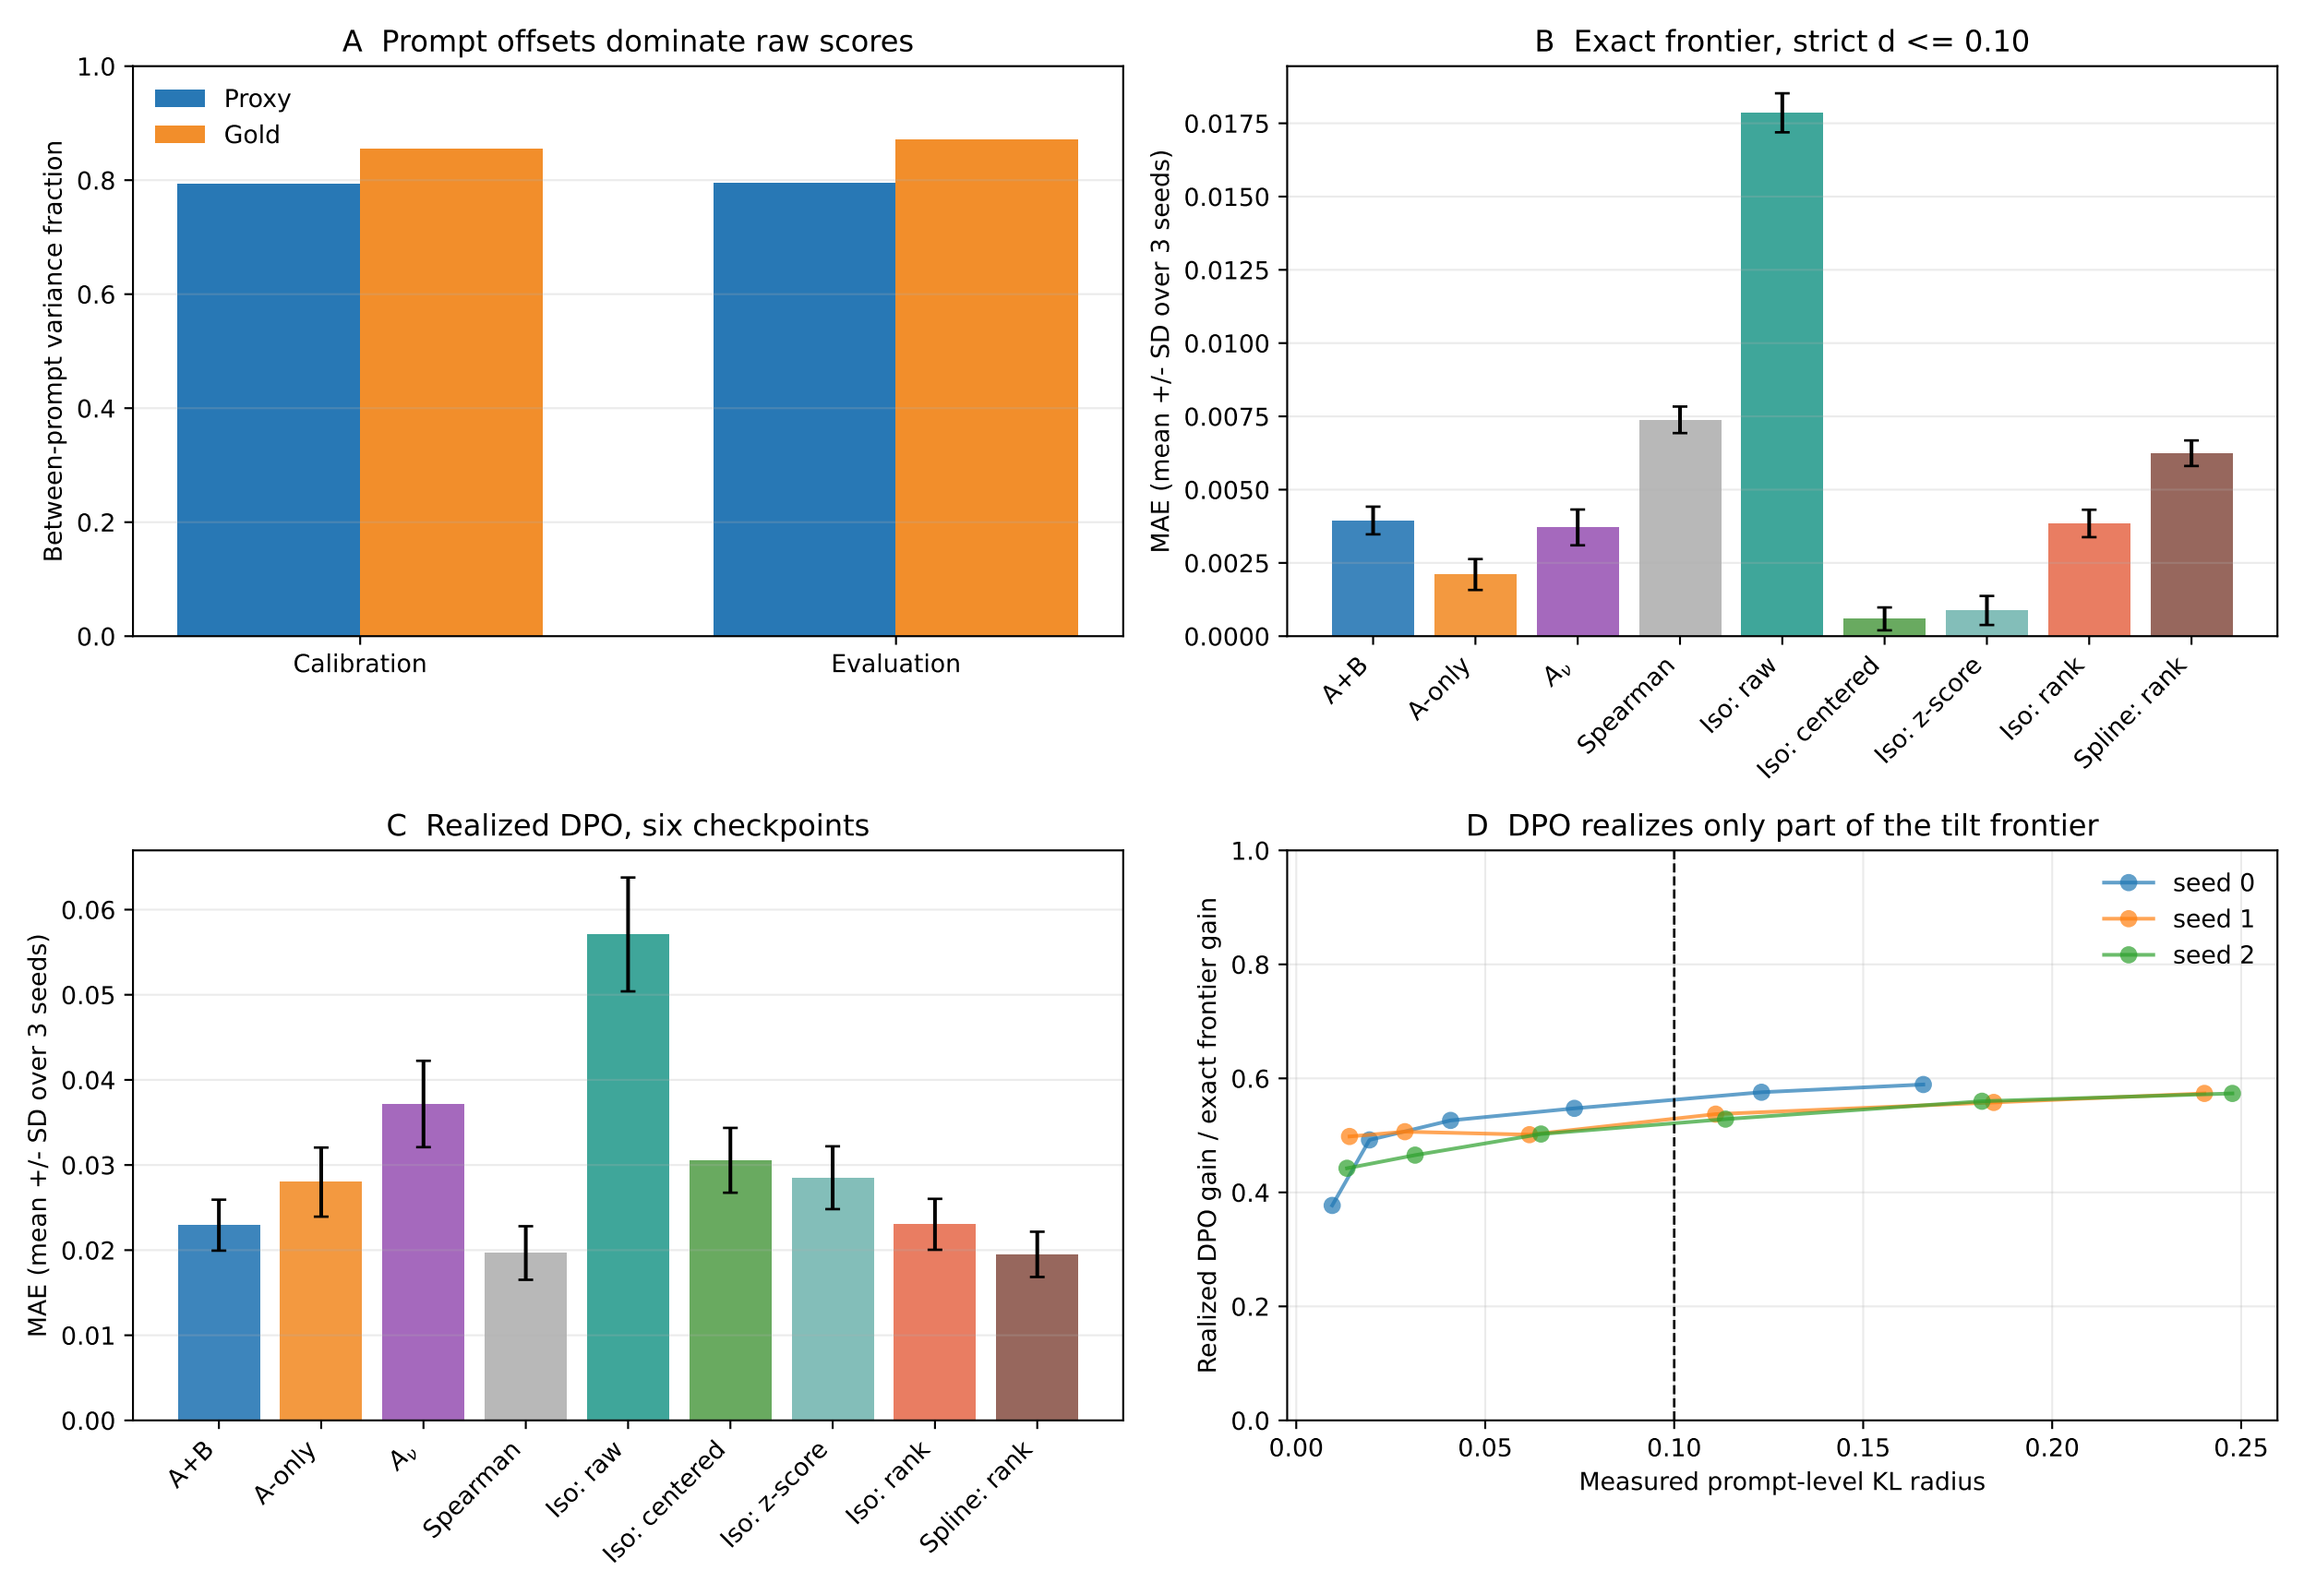
A Prompt offsets dominate raw scores
Proxy
Gold
Between-prompt variance fraction
1.0
0.8
0.6
0.4
0.2
0.0
Calibration
Evaluation
B Exact frontier, strict d <= 0.10
MAE (mean +/- SD over 3 seeds)
0.0175
0.0150
0.0125
0.0100
0.0075
0.0050
0.0025
0.0000
A+B
A-only
Aν
Spearman
Iso: raw
Iso: centered
Iso: z-score
Iso: rank
Spline: rank
C Realized DPO, six checkpoints
MAE (mean +/- SD over 3 seeds)
0.06
0.05
0.04
0.03
0.02
0.01
0.00
A+B
A-only
Aν
Spearman
Iso: raw
Iso: centered
Iso: z-score
Iso: rank
Spline: rank
D DPO realizes only part of the tilt frontier
Realized DPO gain / exact frontier gain
1.0
0.8
0.6
0.4
0.2
0.0
seed 0
seed 1
seed 2
0.00
0.05
0.10
0.15
0.20
0.25
Measured prompt-level KL radius